

Parallelization
and R

Parallelization

When and
How?

MapReduce

R: `parallel`
and `foreach`

Rdsm

R on TACC

References

Parallelization and R

Dennis Wylie, UT Bioinformatics Consulting Group

November 19, 2015

Outline

Parallelization
and R

Parallelization

When and
How?

MapReduce

R: `parallel`
and `foreach`

Rdsm

R on TACC

References

- 1 Parallelization
- 2 When and How?
- 3 MapReduce
- 4 R: `parallel` and `foreach`
- 5 Rdsm
- 6 R on TACC

Serial Computing

Parallelization
and R

Parallelization

When and
How?

MapReduce

R: parallel
and foreach

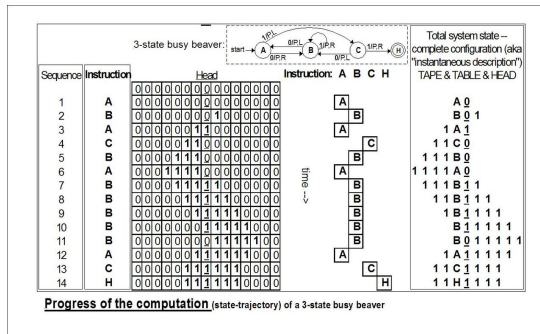
Rdsm

R on TACC

References

What is computation?

- ▶ Many proposed definitions
- ▶ Perhaps most accepted relate to (in many ways equivalent) concepts of **Turing machine** and **lambda calculus**



- ▶ Inherent to the definition of a Turing machine is sequential operation
- ▶ Time assumed to progress in discrete, totally ordered steps

Multiple machines operating simultaneously

- ▶ For some problems, machines may work totally independently
- ▶ May share memory (i.e., same tape)
 - ▶ In practice, share some but not all memory
- ▶ Alternatively, may exchange information (message-passing)
 - ▶ Often one machine passes out work to other machines, waits for them to report back results, then aggregates for further processing

Types of Parallel Machines

Parallelization
and R

Parallelization

When and
How?

MapReduce

R: `parallel`
and `foreach`

`Rdsm`

R on TACC

References

- ▶ Single multicore machine
 - ▶ most modern desktops / laptops
 - ▶ easier to share memory
 - ▶ R: if linux or mac, can use `mclapply`
- ▶ Clusters of multiple machines networked together
 - ▶ harder to share memory
 - ▶ R: can use `clusterApply` family of functions
- ▶ Specialized parallel devices (ignored here)

How Much Faster?

Parallelization
and R

Parallelization

When and
How?

MapReduce

R: parallel
and foreach

Rdsm

R on TACC

References

Generally not all parts of a computation can be parallelized.

Amdahl's Law: if fraction p of the work to be done can be parallelized, the speedup with n processors is

$$\begin{aligned} S &= \frac{(1-p) + p}{(1-p) + \frac{p}{n}} \\ &= \frac{1}{1-p + \frac{p}{n}} \\ &\rightarrow \frac{1}{1-p} \text{ when } n \rightarrow \infty \end{aligned}$$

How Much Faster?

Parallelization
and R

Parallelization

When and
How?

MapReduce

R: parallel
and foreach

Rdsm

R on TACC

References

Generally not all parts of a computation can be parallelized.

Amdahl's Law: if fraction p of the work to be done can be parallelized, the speedup with n processors is

$$\begin{aligned} S &= \frac{(1 - p) + p}{(1 - p) + \frac{p}{n}} \\ &= \frac{1}{1 - p + \frac{p}{n}} \\ &\rightarrow \frac{1}{1 - p} \text{ when } n \rightarrow \infty \end{aligned}$$

But this doesn't take into account the extra overhead incurred if processors have to communicate with each other...

Hill & Marty (2008) provides some (relatively) recent commentary on Amdahl's law.

Techniques for Parallelization

Parallelization
and R

Parallelization

When and
How?

MapReduce

R: parallel
and foreach

Rdsm

R on TACC

References

From Blelloch & Maggs (2004):

- ▶ **Divide and Conquer** split problem into independent subproblems, solve, aggregate results (main approach we will discuss here)
- ▶ **Randomization** useful when optimal decomposition difficult to determine
 - ▶ **Sampling** e.g., to sort a list, select a few elements at random, sort these, use as interval boundaries
 - ▶ **Symmetry Breaking** to get started dividing up a problem
 - ▶ **Load Balancing** divide elements of list into random sublists instead of first n , second n , etc.
- ▶ **Parallel Pointer Techniques** useful for replacing traditional sequential algorithms on lists, trees, and graphs
 - ▶ **Pointer Jumping**
 - ▶ **Euler Tour** replace undirected tree with directed graph, construct linked structure representing Euler tour; good for tree traversal
 - ▶ **Graph Contraction**
 - ▶ **Ear Decomposition** partition of graph edges into ordered collection of paths; first is cycle, others "ears" with endpoints anchored on previous paths
- ▶ **Others**

Example: Pointer Jumping for Tree Root Finding

Parallelization
and R

Parallelization

When and
How?


MapReduce

R: parallel
and foreach


Rdsm

R on TACC


References



Initial Forest



First Iteration



Second Iteration

from https://en.wikipedia.org/wiki/Pointer_jumping

Minor Digression: Functional Programming

Parallelization
and R

The lambda calculus of Alonzo Church emphasizes computation as the application of functions to map inputs into outputs.

Parallelization

When and
How?

Functional programming languages like Lisp (Sussman & Abelson (1983)) implement the ideas of the lambda calculus:

MapReduce

R: `parallel`
and `foreach`

First class functions are treated just like other objects (numbers, strings, etc.); can be passed to other functions as arguments or returned as values

Rdsm

R on TACC

References

Immutability output value of a function depends only on the values of the arguments passed to the function; repeated function calls with same arguments always give same output.

Minor Digression: Functional Programming

Parallelization
and R

The lambda calculus of Alonzo Church emphasizes computation as the application of functions to map inputs into outputs.

Parallelization

When and
How?

Functional programming languages like Lisp (Sussman & Abelson (1983)) implement the ideas of the lambda calculus:

MapReduce

R: parallel
and foreach

First class functions are treated just like other objects (numbers, strings, etc.); can be passed to other functions as arguments or returned as values

Rdsm

R on TACC

References

Immutability output value of a function depends only on the values of the arguments passed to the function; repeated function calls with same arguments always give same output.

Immutability prevents many of the “race conditions” which can plague attempts at parallelization.

Minor Digression: Functional Programming

Parallelization
and R

The lambda calculus of Alonzo Church emphasizes computation as the application of functions to map inputs into outputs.

Parallelization

Functional programming languages like Lisp (Sussman & Abelson (1983)) implement the ideas of the lambda calculus:

When and
How?

MapReduce

First class functions are treated just like other objects (numbers, strings, etc.); can be passed to other functions as arguments or returned as values

R: parallel
and foreach

Rdsm

R on TACC

References

Immutability output value of a function depends only on the values of the arguments passed to the function; repeated function calls with same arguments always give same output.

Immutability prevents many of the “race conditions” which can plague attempts at parallelization.

R supports a functional programming style ...

One of the main programming models for constructing parallel algorithms using functional programming techniques.

Map (more commonly known in R as `(l/m)apply`)

- ▶ $\text{Map}(f, l)$ takes another function f and a list l with entries l_i and
- ▶ returns a new list with entries $f(l_i)$.

Reduce (known in R as `Reduce`)

- ▶ takes a function g and list l of length n ,
- ▶ recursively defines the quantities $r_i = g(r_{i-1}, l_i)$, and
- ▶ returns the single value r_n .

MapReduce

Parallelization
and R

Parallelization

When and
How?

MapReduce

R: parallel
and foreach

Rdsm

R on TACC

References

This can be used to implement the “divide-and-conquer” approach to parallel programming:

- ▶ in the Map step, each computation $f(l_i)$ can be farmed out to a different processor.

Note that both Map and Reduce

- ▶ take functions (f and g) as arguments, and that
- ▶ immutability implies that it doesn't matter what order the processors finish their work in (since each computation depends only on f and l_i).

MapReduce

Parallelization
and R

Parallelization

When and
How?

MapReduce

R: parallel
and foreach

Rdsm

R on TACC

References

This can be used to implement the “divide-and-conquer” approach to parallel programming:

- ▶ in the Map step, each computation $f(l_i)$ can be farmed out to a different processor.

Note that both Map and Reduce

- ▶ take functions (f and g) as arguments, and that
- ▶ immutability implies that it doesn't matter what order the processors finish their work in (since each computation depends only on f and l_i).

Google has invested a lot of effort in further developing the MapReduce model (e.g., see Dean & Ghemawat (2008)) for parallel computing.

MapReduce

Parallelization
and R

Parallelization

When and
How?


MapReduce

R: parallel
and foreach

Rdsm

R on TACC

References



from https://en.wikipedia.org/wiki/File:Mapreduce_Overview.svg

Example: Matrix Multiplication

Parallelization
and R

Matrix multiplication

Parallelization

When and
How?

MapReduce

R: parallel
and foreach

Rdsm

R on TACC

References

$$w_i = \sum_{j=1}^s M_{ij} v_j,$$

where $i \in \{1, \dots, r\}$, is ripe for MapReduce:

Split M into $M^{(b)}$ for $b \in \{1, \dots, B\}$ with $B \leq r$ and each row index i assigned to exactly one block b .

Map step: for each block b (in parallel), calculate

$$w_i^{(b)} = \sum_{j=1}^s M_{ij}^{(b)} v_j \text{ for all } i \text{ assigned to block } b.$$

Reduce step: concatenate/interleave vectors $w^{(b)}$ into single vector w .

mclapply

Parallelization
and R

Parallelization

When and
How?

MapReduce

R: parallel
and foreach

Rdsm

R on TACC

References

```
## serial step to separate M into blocks
Mblocks = list()
for (i in 1:nPieces) {
  lower = (i-1) * rFrac + 1
  upper = ifelse(i==nPieces, r, i*rFrac)
  Mblocks[[i]] = M[lower:upper, ]
}

## parallel step to map the multiplication across blocks
wPieces = mclapply(
  X = Mblocks,
  FUN = function(Mb) {Mb %*slow% v},
  mc.cores = 3
)

## serial step to reduce the pieces into aggregated whole
wmc = Reduce(f=c, x=wPieces)
```

foreach

Parallelization
and R

Parallelization

When and
How?

MapReduce

R: parallel
and foreach

Rdsm

R on TACC

References

```
## parallel step to map the multiplication across blocks
wPieces = foreach (i=1:nPieces) %dopar% {
  ## note syntax:
  ## (1) (i=...) instead of (i in ...)
  ## (2) %dopar%
  lower = (i-1) * rFrac + 1
  upper = ifelse(i==nPieces, r, i*rFrac)
  return(M[lower:upper, ] %*slow% v)
}
```

```
## serial step to reduce the pieces into aggregated whole
wmc = Reduce(f=c, x=wPieces)
```

clusterApply

Parallelization
and R

```
cl = makeCluster(3) ## how many cores to use
```

Parallelization

```
## serial step to separate M into blocks
```

```
Mblocks = list()
```

When and
How?

```
for (i in 1:nPieces) {
```

```
  lower = (i-1) * rFrac + 1
```

```
  upper = ifelse(i==nPieces, r, i*rFrac)
```

```
  Mblocks[[i]] = M[lower:upper, ]
```

```
}
```

MapReduce

R: parallel
and foreach

Rdsm

R on TACC

References

```
## cluster nodes needs to know about `"%*slow%"` and v
```

```
clusterExport(cl, "%*slow%")
```

```
clusterExport(cl, "v")
```

```
## parallel step to map the multiplication across blocks
```

```
wPieces = clusterApply(
```

```
  cl = cl,
```

```
  x = Mblocks,
```

```
  fun = function(x) {x %*%slow% v}
```

```
)
```

```
## serial step to reduce the pieces into aggregated whole
```

```
wcl = Reduce(f=c, x=wPieces)
```

Multithreaded Programming

Parallelization
and R

Parallelization

When and
How?

MapReduce

R: parallel
and foreach

Rdsm

R on TACC

References

Threads are sequences of computation which can be executed concurrently but which can also share resources (memory, value of variables), in contrast with **processes**.

Multiple threads may be components of one process.

The package **Rdsm** by Norm Matloff builds on parallel package to support threaded programming for R:

`mgrmakevar` allows creation of shared variables

`myinfo$id` provides id of thread currently executing

`rdsmlock` provides mutex lock

`barr` provides barrier to keep threads in sync

Multithreaded Programming

Parallelization
and R

Parallelization

When and
How?

MapReduce

R: parallel
and foreach

Rdsm

R on TACC

References

Threads are sequences of computation which can be executed concurrently but which can also share resources (memory, value of variables), in contrast with **processes**.

Multiple threads may be components of one process.

The package **Rdsm** by Norm Matloff builds on parallel package to support threaded programming for R:

`mgrmakevar` allows creation of shared variables

`myinfo$id` provides id of thread currently executing

`rdsmlock` provides mutex lock

`barr` provides barrier to keep threads in sync

Rdsm (and parallel computing in “data science” in general) is treated well in Matloff (2013).

Race Conditions

Parallelization
and R

Parallelization

When and
How?

MapReduce

R: parallel
and foreach

Rdsm

R on TACC

References

When multiple threads concurrently execute while sharing mutable state variables, race conditions can occur:

```
## unreliable function
s = function(n) {
  for (i in 1:n) {
    tot[1, 1] = tot[1, 1] + 1
  }
}

c2 = makeCluster(2)
clusterExport(c2, "s")
mgrinit(c2)
mgrmakevar(c2, "tot", 1, 1)
tot[1, 1] = 0
clusterEvalQ(c2, s(1000))
tot[1, 1]  ## should be 2000, but likely far from it
```

Mutual Exclusion (Mutex/Lock)

Parallelization
and R

Parallelization

When and
How?

MapReduce

R: parallel
and foreach

Rdsm

R on TACC

References

```
s1 = function(n) {  
  require(Rdsm)  
  for (i in 1:n) {  
    rdsmlock("totlock")  
    tot[1, 1] = tot[1, 1] + 1  
    rdsmunlock("totlock")  
  }  
}
```

```
c2 = makeCluster(2)  
clusterExport(c2, "s1")  
mgrinit(c2)  
mgrmakevar(c2, "tot", 1, 1)  
tot[1, 1] = 0  
mgrmakelock(c2, "totlock")  
clusterEvalQ(c2, s1(1000))
```


Matrix Multiplication: Rdsm

Parallelization
and R

Parallelization

When and
How?

MapReduce

R: parallel
and foreach

Rdsm

R on TACC

References

```
slowMultThread = function(w, M, v) {  
  require(parallel)  
  
  ## determine which rows this thread will handle  
  idx = splitIndices(nrow(M), myinfo$nrwrks)[[myinfo$id]]  
  w[idx, ] = M[idx, ] %*%slow% v[ , ]  
  
  return(0)  ## don't do expensive return of result  
}  
  
cl = makeCluster(3); mgrinit(cl)  
  
## cluster needs to define its (shared) versions of w, M, v  
mgrmakevar(cl, "Mcl", nrow(M), ncol(M))  
Mcl[ , ] = M  
mgrmakevar(cl, "vcl", length(v), 1)  
vcl[ , ] = v  
mgrmakevar(cl, "wcl", nrow(M), 1)  
  
## cluster needs to know about slowMult and friends  
clusterExport(cl, c("slowMult", "%*slow%", "slowMultThread"))  
  
## parallel step to do the multiplication  
clusterEvalQ(cl, slowMultThread(wcl, Mcl, vcl))
```

Pointer Jumping Redux

Parallelization
and R

Parallelization

When and
How?


MapReduce

R: parallel
and foreach


Rdsm

R on TACC


References



Initial Forest



First Iteration



Second Iteration

Important during pointer jumping to keep jump steps in sync.

Why? Imagine result if processor for node 11 didn't finish first iteration until after all others were done with both. . .

References

Why? Imagine result if processor for node 11 didn't finish first iteration until after all others were done with both. . . then 12 would end up pointing to 9, not to 5.

Pointer Jumping Redux

Parallelization
and R

Parallelization

When and
How?


MapReduce

R: parallel
and foreach


Rdsm

R on TACC


References



Initial Forest



First Iteration



Second Iteration

Important during pointer jumping to keep jump steps in sync.

Why? Imagine result if processor for node 11 didn't finish first iteration until after all others were done with both. . . then 12 would end up pointing to 9, not to 5.

One method for enforcing synchronization is to use **barrier**.

Pointer Jumping and Barriers

Parallelization
and R

Parallelization

When and
How?

MapReduce

R: parallel
and foreach

Rdsm

R on TACC

References

```
iterParBar = function(nwpntcl, pntcl, ncl) {  
  require(parallel)  
  
  ## determine which rows this thread will handle  
  idx = splitIndices(length(pntcl), myinfo$nrwrks)[[myinfo$id]]  
  
  for (i in 1:ncl[1, 1]) {  
    ## do the jumping, store results in new variable  
    nwpntcl[idx] = pntcl[ pntcl[idx] ]  
  
    ## wait for all threads to finish jumping  
    barr()  
  
    ## reset old pointer to new, jumped values  
    pntcl[idx] = nwpntcl[idx]  
  
    ## wait for all threads to finish resetting  
    barr()  
  }  
  
  ## don't do expensive return of result  
  return(0)  
}
```

Dining Philosophers

Parallelization
and R

Parallelization

When and
How?

MapReduce

R: parallel
and foreach

Rdsm

R on TACC

References



TACC provides a bunch of multicore machines clustered together ready for you to use these techniques on. . .

As for how to do it, you've got a few options:

- ▶ Interactive Session:
 - ▶ Use RStudio on the TACC Visualization Portal (<https://vis.tacc.utexas.edu/>)
 - ▶ if you have your own version of R installed on TACC, this may not work!
 - ▶ ssh in, start an idev session and run R on the cores you've got available there

R on TACC: Batch Processing

Parallelization
and R

Parallelization

When and
How?

MapReduce

R: parallel
and foreach

Rdsm

R on TACC

References

TACC provides a bunch of multicore machines clustered together ready for you to use these techniques on...

As for how to do it, you've got a few options:

- ▶ Batch Processing (using slurm):
 - ▶ Write:
 - ▶ R script file(s) you want to run
 - ▶ job file containing calls to Rscript
 - ▶ Generate .slurm file (e.g., using `launcher_creator.py` from BiolTeam bin)
 - ▶ Submit job: `sbatch <job>.slurm`

References I

Parallelization and R

Parallelization

When and How?

MapReduce

R: parallel and foreach

Rdsm

R on TACC

References

Blelloch, Guy E, & Maggs, Bruce M. 2004. *Parallel Algorithms*.

Dean, Jeffrey, & Ghemawat, Sanjay. 2008. MapReduce: Simplified Data Processing on Large Clusters. *Communications of the ACM*, **51**(1), 107–113.

Hill, Mark D, & Marty, Michael R. 2008. Amdahl's Law in the Multicore Era. *Computer*, 33–38.

Matloff, Norman. 2013. *Parallel Computing for Data Science*. University of California, Davis.

Sussman, Gerry, & Abelson, Harold. 1983. *Structure and Interpretation of Computer Programs*. MIT Press, Cambridge, Mass.